

DATA MINING

DIGITAL ASSIGNMENT - 05

ABHIRUPA – 17BCE0437

Clustering

- Implement a k-means algorithm with appropriate package to partition observations in a dataset into a specific number of clusters in order to aid in analysis of the data.
 - Use Toolkit / Package to perform the process
 - Devise an elbow curve to select the optimal number of clusters (k)
 - Generate and visualise a k-means clustering algorithm

Note : Dataset in CSV can be generated or downloaded from the internet. Please specify the source of the dataset in the documentation steps of this program.

DATA SET USED: Titanic dataset from Kaggle

DATA SET LINK:

<http://s3.amazonaws.com/assets.datacamp.com/course/Kaggle/train.csv>

```
In [1]: 1 import pandas as pd
2 import numpy as np
3 from sklearn.cluster import KMeans
4 from sklearn.preprocessing import LabelEncoder
5 from sklearn.preprocessing import MinMaxScaler
6 import seaborn as sns
7 import matplotlib.pyplot as plt
8
```

```
In [2]: 1 data=pd.read_csv("http://s3.amazonaws.com/assets.datacamp.com/course/Kaggle/train.csv")
```

```
In [4]: 1 from sklearn.model_selection import train_test_split
2 train, test = train_test_split(data, test_size = 0.3, random_state = 100)
```

```
In [6]: 1 print(train.columns.values)

['PassengerId' 'Survived' 'Pclass' 'Name' 'Sex' 'Age' 'SibSp' 'Parch'
'Ticket' 'Fare' 'Cabin' 'Embarked']
```

```
In [7]: 1 print("Missing values in training set:")
2 print(train.isna().sum())
3 print("\n")
4 print("Missing values in testing set:")
5 print(test.isna().sum())
```

```
Missing values in training set:
PassengerId    0
Survived        0
Pclass          0
Name            0
Sex             0
Age            128
SibSp           0
Parch           0
Ticket          0
Fare            0
Cabin          473
Embarked        1
dtype: int64
```

```
Missing values in testing set:
PassengerId    0
Survived        0
Pclass          0
Name            0
Sex             0
Age            49
SibSp           0
Parch           0
Ticket          0
Fare            0
Cabin          214
Embarked        1
dtype: int64
```

```
In [11]: 1 # Replacing missing values by mean imputation, ie, with mean column values
2 train.fillna(train.mean(), inplace=True)
3 test.fillna(test.mean(), inplace=True)
```

/home/abhirupa/anaconda3/lib/python3.7/site-packages/pandas/core/generic.py:5434: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>
self._update_inplace(new_data)

```
In [12]: 1 print("Missing values in training set:")
2 print(train.isna().sum())
3 print("\n")
4 print("Missing values in testing set:")
5 print(test.isna().sum())
```

```
Missing values in training set:
PassengerId    0
Survived        0
Pclass          0
Name            0
Sex             0
Age             0
SibSp           0
Parch           0
Ticket          0
Fare            0
Cabin          473
Embarked        1
dtype: int64
```

```

Missing values in testing set:
PassengerId      0
Survived          0
Pclass           0
Name             0
Sex              0
Age              0
SibSp            0
Parch            0
Ticket           0
Fare             0
Cabin            214
Embarked         1
dtype: int64

```

```

In [13]: 1 # Dropping non-numeric data fields
          2 train = train.drop(['Name', 'Ticket', 'Cabin', 'Embarked'], axis=1)
          3 test = test.drop(['Name', 'Ticket', 'Cabin', 'Embarked'], axis=1)

```

```


In [14]: 1 # Use label encoding to convert 'Sex' feature into numeric format
          2 labelEncoder = LabelEncoder()
          3 labelEncoder.fit(train['Sex'])
          4 labelEncoder.fit(test['Sex'])
          5 train['Sex'] = labelEncoder.transform(train['Sex'])
          6 test['Sex'] = labelEncoder.transform(test['Sex'])
          7

```

```

In [22]: 1 mms = MinMaxScaler()
          2 mms.fit(train)
          3 data_transformed = mms.transform(train)
          4
          5 Sum_of_squared_distances = []
          6 K = range(1,15)
          7 for k in K:
          8     km = KMeans(n_clusters=k)
          9     km = km.fit(data_transformed)
         10     Sum_of_squared_distances.append(km.inertia_)
         11
         12 plt.plot(K, Sum_of_squared_distances, 'bx-')
         13 plt.xlabel('k')
         14 plt.ylabel('Sum_of_squared_distances')
         15 plt.title('Elbow Method For Optimal k')
         16 plt.show()
         17
         18 # k=5

```



```

In [59]: 1 from sklearn.metrics import accuracy_score
          2 # Dropping the feature that is to be tested
          3 X = np.array(train.drop(['Survived'], 1).astype(float))
          4 y = np.array(train['Survived'])
          5 xtest= np.array(test.drop(['Survived'], 1).astype(float))
          6 ytest= np.array(test['Survived'])
          7

```

```

In [58]: 1 kmeans = KMeans(n_clusters=2) # You want cluster the passenger records into 2: Survived or Not survived
          2 kmeans.fit(X)
          3 print(kmeans.cluster_centers_)
          4 print(kmeans.labels_)
          5 correct = 0
          6 for i in range(len(X)):
          7     predict_me = np.array(X[i].astype(float))
          8     predict_me = predict_me.reshape(-1, len(predict_me))
          9     prediction = kmeans.predict(predict_me)
         10     if prediction[0] == y[i]:
         11         correct += 1
         12
         13 print("ACCURACY:", correct/len(X))

```

```
[[6.73725552e+02 2.28075710e+00 6.84542587e-01 3.12424338e+01  
4.73186120e-01 3.31230284e-01 3.16731726e+01]  
[2.24055556e+02 2.30065359e+00 6.07843137e-01 2.88393791e+01  
6.14379085e-01 4.11764706e-01 3.39641464e+01]]  
[1 1 0 1 0 0 1 0 1 0 1 0 1 1 1 1 1 0 1 0 1 0 0 1 1 1 0 0 0 0 1 0 0 0 0 1 0  
0 0 0 1 0 0 1 0 1 1 0 0 0 0 0 0 0 1 1 1 1 1 1 1 0 1 0 0 0 1 1 0 1 0 1 1  
0 0 1 0 1 0 0 0 0 1 1 1 0 0 1 1 1 1 1 0 0 0 0 1 0 0 1 0 0 0 1 0 0 0 1 0 1  
1 0 0 0 0 0 0 1 0 0 0 0 1 0 0 0 1 1 0 1 0 1 1 0 0 1 0 1 1 0 0 0 1 0 0 0 1  
1 0 0 1 0 1 1 1 1 0 0 0 0 0 1 0 1 0 1 1 1 1 1 0 1 0 1 0 1 1 0 1 1 0 0 0  
1 1 1 1 1 1 0 1 1 0 1 0 0 1 0 1 1 0 1 1 1 0 0 1 1 0 0 1 0 0 0 1 1 1 0 1 1  
1 0 0 0 0 1 1 0 0 0 1 0 0 1 1 0 1 0 1 0 1 1 0 1 0 1 0 0 1 0 1 0 1 1 0 0 1  
0 1 0 0 1 0 1 0 0 0 1 0 0 0 1 0 1 1 0 0 0 0 1 1 1 1 1 0 0 1 1 1 0 0 0 0 1  
1 1 1 1 1 1 1 1 1 0 0 1 0 1 0 0 1 0 0 1 0 1 0 1 1 1 0 0 0 1 0 0 1 1 0 0  
0 1 0 1 0 0 0 0 0 1 0 0 0 1 0 0 1 1 0 0 0 0 1 1 0 1 1 1 1 0 0 0 1 0 0 1  
0 0 1 1 1 1 0 0 0 1 0 0 1 1 0 0 1 0 1 0 0 1 1 1 1 0 0 0 0 0 1 0 1 1 0 1 1  
0 0 1 1 1 0 0 1 0 1 1 1 1 0 1 0 0 0 1 1 0 0 1 0 1 0 1 0 0 1 1 0 0 0 1 0 1  
1 1 0 1 0 0 0 0 1 0 0 1 1 1 1 1 0 1 0 0 1 1 1 1 0 0 0 1 1 1 0 0 1 0 1 0 0  
1 1 1 1 0 0 0 1 1 1 0 0 1 1 0 1 0 0 1 1 0 1 0 1 1 1 0 0 0 1 1 0 1 1 1 0 0  
1 1 1 1 1 0 0 0 1 1 0 1 1 1 0 1 1 1 0 0 0 1 1 0 1 0 0 1 1 0 1 1 1 0 0 1 0  
1 0 1 1 0 0 0 1 0 1 1 1 1 0 0 1 1 1 0 1 1 1 1 0 0 1 1 0 0 0 1 0 1 1 1 0 0  
1 0 0 0 1 1 0 1 1 1 0 0 0 1 1 0 1 0 0 0 0 0 1 0 1 1 0 0 0 0 0]
```

ACCURACY: 0.5296950240770465

In []: 1